

## Report

- Code with Detailed Explanation

- Calculate kernel function

$$k_{RQ}(x_i, x_j) = \sigma^2 \left( 1 + \frac{(x_i - x_j)^2}{2\alpha l^2} \right)^{-\alpha}$$

Where  $(x_i - x_j)^2$  is RQ distance,  $l$  determines the length of wiggles in the function,  $\sigma$  is the average distance of the function away from its mean,  $\alpha$  determines the relative weighting of large scale and small scale variations. When  $\alpha \rightarrow \infty$ , rational quadratic kernel will be identical to standard exponential kernel (a.k.a Gaussian kernel).

```
def kernel(X1, X2, a, l, sigma):  
    '''  
    X1: m points  
    X2: n points  
    K :(m,n)  
    '''  
  
    rqdist2 = np.sum(X1**2, 1).reshape(-1, 1) + np.sum(X2**2, 1) - 2 * X1 @ X2.T  
    return sigma**2 * (1 + rqdist2/(2*a*l**2))**(-a)
```

- Derive negative marginal log-likelihood

$y_n = f(x_n) + \epsilon_n$  , where  $\epsilon \sim N(\cdot | 0, \beta^{-1})$  and  $\beta = 5$

$$p(y|f) = N(y|f, \beta^{-1}I_N)$$

$$p(f) = N(0, K)$$

$$P(y) = \int p(y|f)p(f)df = N(y|0, C_N) = \frac{1}{\sqrt{(2\pi)^N} |C_N|^{1/2}} \exp\left(-\frac{1}{2}y^T C_N^{-1}y\right)$$

where  $C_N(x_i, x_j) = k(x_i, x_j) + \beta^{-1}\delta_{ij}$

$$\Rightarrow nll = \frac{1}{2} \log(|C_N|) + \frac{N}{2} \log(2\pi) + \frac{1}{2}y^T C_N^{-1}y$$

```
def nll(X_train, Y_train, noise):  
    Y_train = Y_train.ravel()  
    def nll_naive(theta):  
        C = kernel(X_train, X_train, a=theta[0], l=theta[1], sigma=theta[2]) + \\\n        noise*np.identity(N) # noise = 1/beta = 0.2  
        return 0.5*np.log(det(C)) + 0.5*N*np.log(2*np.pi) + 0.5*Y_train.T @ inv(C) @ Y_train  
    return nll_naive
```

- Optimize kernel hyperparameters  $\theta = (\alpha, l, \sigma)$

Set initial guess  $\theta^{(0)} = (\alpha^{(0)}, l^{(0)}, \sigma^{(0)}) = (1, 1, 1)$  , since all parameters are positive, let's

set their bounds from  $1e-5$  to  $\infty$ , and apply L-BFGS-B method to optimize  $\theta$ .

```
from scipy.optimize import minimize
theta = minimize(nll(X_train, Y_train, noise), x0=[1, 1, 1], method='L-BFGS-B',
                 bounds=((1e-5, None),(1e-5, None),(1e-5, None)))
a, l, sigma = theta.x
```

#### ○ Prediction

Given test data, return their means and variances (in covariance.diagonal()).

$$\mu(x_{N+1}) = k^T C_N^{-1} y, \sigma^2(x_{N+1}) = c + k^T C_N^{-1} k, C_{N+1} = \begin{bmatrix} C_N & k \\ k^T & c \end{bmatrix}$$

```
def posterior(X_test, X_train, Y_train, noise, a, l, sigma):
    C_N = kernel(X_train, X_train, a, l, sigma) + noise * np.identity(N)
    C_N_inv = inv(C_N)
    K = kernel(X_train, X_test, a, l, sigma) # (N,N_test)
    c = kernel(X_test, X_test, a, l, sigma) + noise * np.identity(X_test.shape[0])
                                                    # (N_test, N_test)

    mu_new = K.T @ C_N_inv @ Y_train # (N_test,1)
    cov_new = c - K.T @ C_N_inv @ K # (N_test, N_test)
    return mu_new, cov_new
```

#### ○ Plot

Calculate 95% confidence interval, and plot means with 95% confidence interval.

$$95\% \text{ confidence interval} = \left[ \mu - 1.96 \frac{\sigma}{\sqrt{N}}, \mu + 1.96 \frac{\sigma}{\sqrt{N}} \right]$$

```
def plot(X_train, Y_train, noise, a, l, sigma):
    fig, (ax) = plt.subplots(1, 1, figsize=(8, 6))
    ax.set_ylim(-6, 6)

    # Show all training data points
    ax.scatter(X_train, Y_train, color='steelblue', marker='.')

    X = np.linspace(-60,60,1000)
    mu, cov = posterior(X.reshape(-1,1), X_train, Y_train, noise, a, l, sigma)
    mu, var = mu.ravel(), cov.diagonal()

    # Draw a line to represent the mean of f in range
    ax.plot(X, mu, color='steelblue')

    # Mark 95% confidence interval of f
    confidence = 1.96*(np.sqrt(var/34))
    ax.fill_between(X, mu-confidence, mu+confidence, color='steelblue', alpha=0.5)
    plt.show()
```

- Experiments Settings and results

Set distribution of noise:  $y_n = f(x_n) + \epsilon_n$ , where  $\epsilon \sim N(\cdot | 0, \beta^{-1})$  and  $\beta = 5$


Set initial guess at  $\theta = (1, 1, 1)$ , with each bound from  $1e-5$  to  $\infty$ , and then apply L-BFGS-B method to optimize  $\theta$ .

We get optimized  $\theta = (\alpha, l, \sigma) = (275.5275, 3.3116, 1.3107)$


with negative log-likelihood = 50.6827.

By prediction of each  $x$  in  $[-60, 60]$ , and plot with mean and 95% confidence interval.

(Code explanations are on preceding page)



Set initial guess at  $\theta = (100, 100, 100)$ , and we get optimized  $\theta = (\alpha, l, \sigma) = (100.7633, 100.6455, 99.9756)$ , negative log-likelihood = 153.2907




It's clear that since L-BFGS-B method is gradient-based method, the result is vulnerable to starter value/initial guess. If the initial guess is far away from the global optimum, then it may be easy to get stuck in a local optimum. Thus, we can generate a population of initial guesses to relieve this problem.

- Observations and Discussion

$l$  determines the length of wiggles in the function,  $\sigma$  is the average distance of the function away from its mean,  $\alpha$  determines the relative weighting of large scale and small scale variations. (By hand-tuning hyperparameters, we can see how each hyperparameter affects our model)


It's clear to observe how  $l$  and  $\sigma$  affect our model in the cases below.




As  $\alpha$  increases, the prediction function becomes steeper; as  $l$  increases, the length of wiggles become larger, the density of peaks declines, and the prediction function becomes smoother; as  $\sigma$  increases, the variations becomes larger (i.e. further away from mean).

Let's take a further look at  $\alpha$ :


$(\alpha, l, \sigma) = (1, 1, 1)$  , nll = 55.9227




$(\alpha, l, \sigma) = (0.5, 1, 1)$  , nll = 56.8214



$(\alpha, l, \sigma) = (0.1, 1, 1)$  , nll = 66.8600



$(\alpha, l, \sigma) = (0.01, 1, 1)$  , nll = 99.4035



It's clear that as  $\alpha$  decreases, the prediction function becomes biased/less accurate, yet the number of peaks are the same.